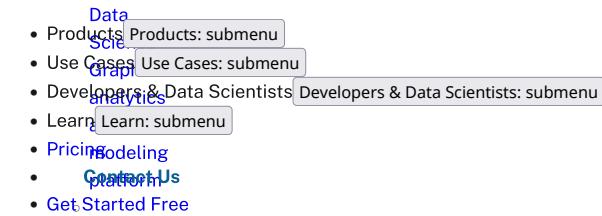


[Skip to content](#)The GraphRAG Manifesto: Unlock Better GenAI Results With Knowledge Graphs | [Read Now](#)

[Get Started](#)
[Contact Us](#) [Back](#)

- [Aura Login](#)
- [Partners](#) Partners: submenu
 - [Find a Partner](#)
 - [Become a Partner](#)
 - [Solution Partners](#)
 - [Managed](#)
 - [OEM Partners](#)
 - [Deploy Anywhere](#)
 - [Partner Portal Login](#)
- [Company](#) Company: submenu
 - [About Us](#)
 - [Newsroom](#)
 - [Awards and Honors](#)
 - [Graphs4Good](#)
 - [Careers](#)
 - [Culture](#)
 - [Diversity](#)
 - [Leadership](#)

• [Support](#)
[Search](#)



(Neo4j Developer Blog) [\[BACK\]](#)
[Center](#)

Entity Linking and Relationship Extraction With Relik in LlamalIndex



Tomaž Bratanić Aug 12 7 mins read

[Contact Us](#)

Neo4i

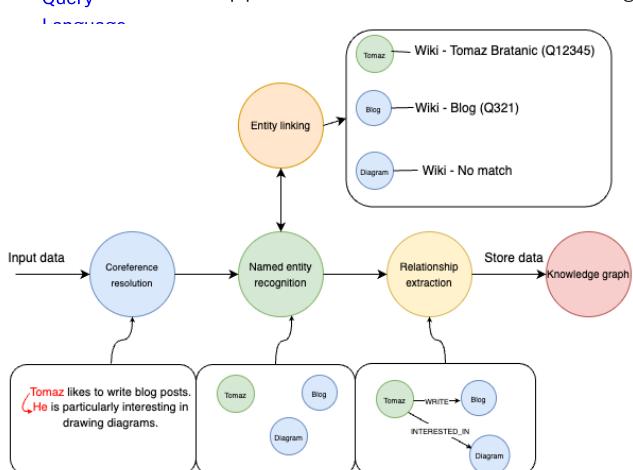


visualization

Constructing knowledge graphs from text has been a fascinating area of research for quite some time. With the advent of large language models (LLMs), this field has gained more mainstream attention. However, LLMs can be quite costly. An alternative approach is to fine-tune smaller models, which has been supported by academic research, yielding more efficient solutions. Today, we will explore [Relik](#), a framework for running blazing fast and lightweight information extraction models, developed by the NLP group at the Sapienza University of Rome.

Cypher

A typical information extraction pipeline without an LLM looks like the following:



Information extraction pipeline — image by author

- Use

The diagram illustrates an information extraction pipeline, starting from input data that consists of text mentioning "Tomaz likes to write blog posts. He is particularly interested in drawing diagrams." The process begins with coreference resolution to identify "Tomaz" and "He" as the same entity. Named entity recognition (NER) then identifies entities such as "Tomaz," "Blog," and "Diagram."

Contact Us

Entity linking is the process that follows NER, where recognized entities are mapped to corresponding entries in a database or knowledge base. For example, “Tomaz” is linked to “Tomaz Bratanic (Q12345)” and “Blog” to “Blog (Q321),” but “Diagram” has no match in the knowledge base.

AI
Relationship extraction is the subsequent step where the system identifies and extracts meaningful relationships between the recognized entities. This example identifies that “Tomaz” has a relationship with “Blog” characterized by “WRITES,” indicating that Tomaz writes blogs. Additionally, it identifies that “Tomaz” has a relationship with “Diagram” characterized by “INTERESTED_IN,” indicating that Tomaz is interested in diagrams.

LLMs
Finally, this structured information, including the entities and their relationships, is stored in a knowledge graph, allowing for organized and accessible data for further analysis or retrieval.

Traditionally, without the power of LLMs, this entire process relies on a suite of specialized models, each handling a specific task from coreference resolution to relationship extraction. While integrating these models demands more effort and coordination, **LLMs** offer a significant advantage: reduced costs. By fine-tuning smaller, task-specific models, the overall expense of building and maintaining the system can be kept in check.

The code is available on [GitHub](#).

Environment Setup

AI
I suggest you use a separate Python environment like [Google Colab](#), as we will have to play around with dependencies a bit. The models are faster on GPU, so you can use a GPU-powered runtime if you have the Pro version.

Industries
Additionally, we need to set up Neo4j, a native graph database, to store the extracted information. There are many ways to [set up your database instance](#). However, I recommend using [Neo4j Aura](#), which provides a free cloud instance that can easily be accessed from a Google Colab notebook.

Use

[Neo4j Aura](#) is a fully managed cloud solution.

Fraud

After the database has been created, we can define a connection using LlamaIndex:

Detection

Knowledge

```
from llama_index.graph_stores.neo4j import Neo4jPGStore

username="neo4j"
password="rubber-cuffs-radiator"
url="bolt://54.89.19.156:7687"

graph_store = Neo4jPGStore(
    username=username,
    password=password,
    url=url,
    refresh_schema=False
)
```

across

Dataset [industries](#)

- We will use a news dataset I obtained via [Diffbot API](#) some time ago. The dataset is conveniently available on GitHub for us to reuse:

```
import pandas as pd

NUMBER_OF_ARTICLES = 100
news = pd.read_csv(
    "https://raw.githubusercontent.com/tomasonjo/blog-datasets/main/news_articles.csv"
)
news = news.head(NUMBER_OF_ARTICLES)
```

Data

Coreference Resolution

- Back

The first step in the pipeline is a coreference resolution model. **Coreference resolution** is the task of identifying all expressions in a text refer to the same entity.

[Developer](#)

Contact Us

To my knowledge, there aren't many open-source models available for coreference resolution. I tried the [maverick-coref](#), but in my tests [Coreferee](#) from spaCy worked better, so we will use that. The only disadvantage of using Coreferee is that we have to deal with dependency hell, which is solved in the notebook, but we'll not go through it here.

You can load the coreference model in spaCy with the following code:

```
import spacy, coreferee

coref_nlp = spacy.load('en_core_web_lg')
coref_nlp.add_pipe('coreferee')
```

The Coreferee model detects clusters of expression that refer to the same entity or entities. To rewrite the text based on these clusters, we have to implement our own function:

```
def coref_text(text):
    coref_doc = coref_nlp(text)
    resolved_text = ""

    for token in coref_doc:
        repres = coref_doc_.coref_chains.resolve(token)
        if repres:
            resolved_text += " " + " ".join([
                t.text
                if t.ent_type_ == ""
                else [e.text for e in coref_doc.ents if t in e][0]
                for t in repres
            ])
    return resolved_text
```

Let's test the function to make sure the models and dependencies are set up properly:

```
print(
    coref_text("Tomaz is so cool. He can solve various Python dependencies and not cry")
)
# Tomaz is so cool . Tomaz can solve various Python dependencies and not cry
```

In this example, the model identified that "Tomaz" and "He" refer to the same entity. Using the coref_text function, we replace "He" with "Tomaz."

Note that the rewriting doesn't always return grammatically correct sentences due to using simple replace logic for entities within the cluster. However, it should be good enough for most scenarios.

Now we apply the coreference resolution to our news dataset and wrap the results as Llamaindex documents:

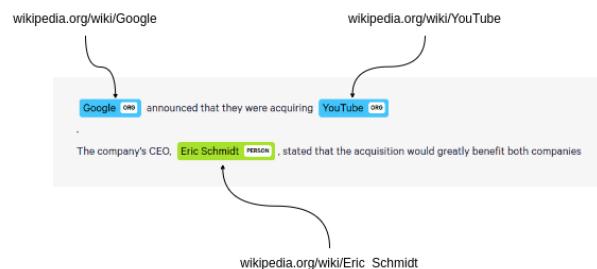
```
from llama_index.core import Document

news["coref_text"] = news["text"].apply(coref_text)
documents = [
    Document(text=f"{row['title']}: {row['coref_text']}") for i, row in news.iterrows()
]
```

Entity Linking and Relationship Extraction

Relik is a library with models for entity linking (EL) and relationship extraction (RE), and it also supports models that combine the two. In entity linking, Wikipedia is used as the target knowledge base to map entities in text to their corresponding entries in the encyclopedia.

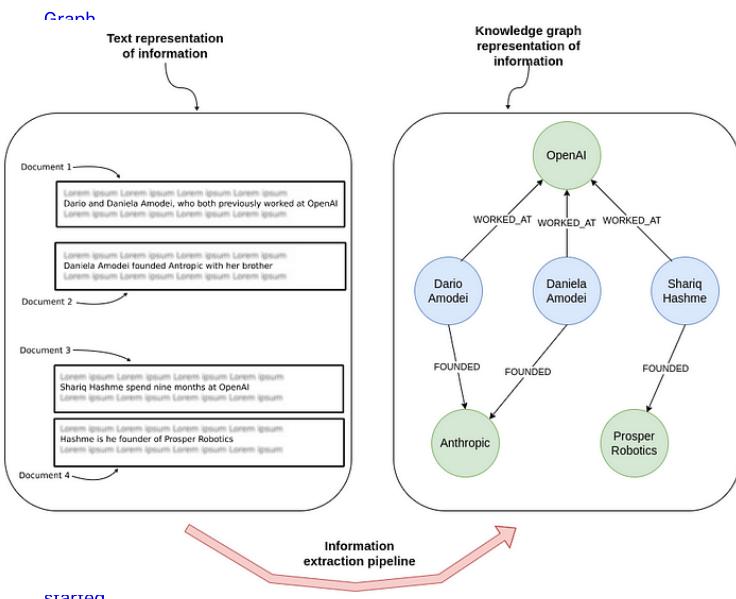
Contact Us



[Science library](#)

Linking entities to Wikipedia – Image by author

On the other hand, relationship extraction involves identifying and categorizing the relationships between entities within a text, enabling the extraction of structured information from unstructured data.



[Relationship extraction – Image by author](#)

If you are using a free Colab version, use the `relik-ie/relik-relation-extraction-small` model, which performs only relationship extraction. If you have a Pro version, or you will use it on a stronger local machine, you can test the `relik-ie/relik-cie-small` model, which performs entity linking and relationship extraction.

```
from llama_index.extractors.relik.base import RelikPathExtractor

relik = RelikPathExtractor(
    model="relik-ie/relik-relation-extraction-small"
)

# Use on Pro Collab with GPU
# relik = RelikPathExtractor(
#     model="relik-ie/relik-cie-small", model_config={"skip_metadata": True, "device": "cuda"}
# )
```

Contact Us

LEARN

```
import os

from llama_index.embeddings.openai import OpenAIEmbedding
from llama_index.llms.openai import OpenAI

os.environ["OPENAI_API_KEY"] = "sk-"

llm = OpenAI(model="gpt-4o", temperature=0.0)
embed_model = OpenAIEmbedding(model_name="text-embedding-3-small")
```

Note that the LLM will not be used during graph construction.

Resource

With everything in place, we can instantiate a PropertyGraphIndex and use the news documents as input data to a knowledge graph.

White

Additionally, we need pass the relik model as the kg_extractors value to extract the relationships:

→

```
from llama_index.core import PropertyGraphIndex

index = PropertyGraphIndex.from_documents(
    documents,
    kg_extractors=[relik],
    llm=llm,
    embed_model=embed_model,
    property_graph_store=graph_store,
    show_progress=True,
)
```

across

After constructing the graph, you can open Neo4j Browser to validate the imported graph. You should get a similar visualization by running the following Cypher statement:

○

```
MATCH p=(:__Entity__)--(:__Entity__)
RETURN p LIMIT 250
```

to

Results know

Graph**Technology**

○ CONNECT

○

Neo4j**Events****Hub****Live**

and

on-

demand

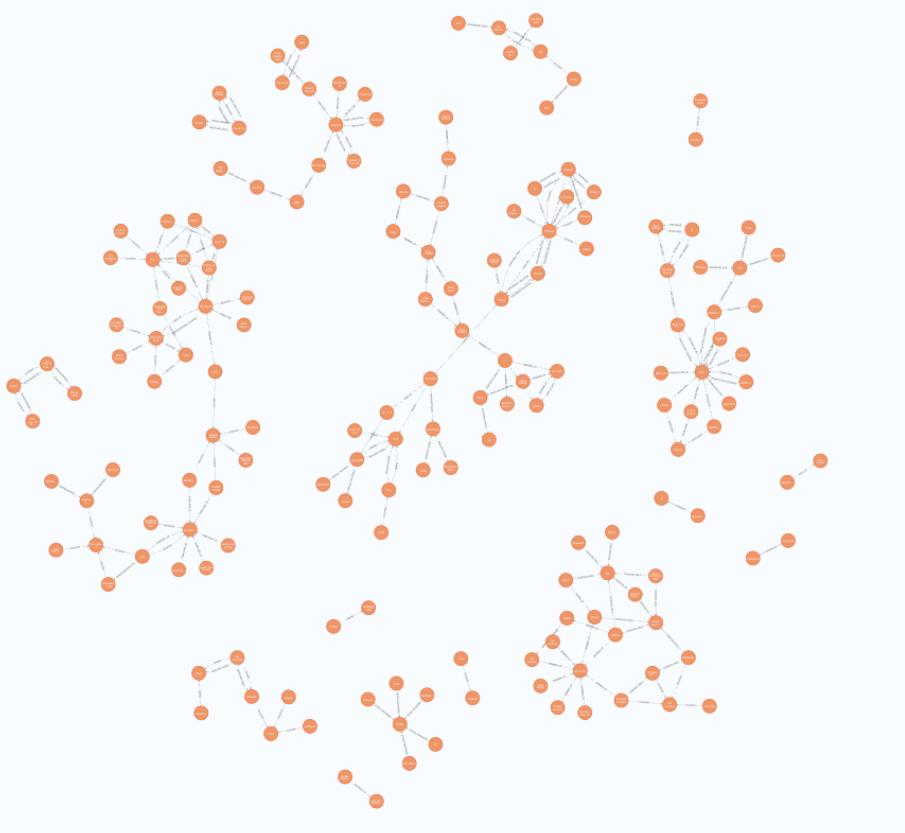
events,

training,

webinars,

and

Contact Us



Graph visualization – Image by author

conference

Question Answering

every

quarter

Using LlamaIndex, it is now easy to perform question answering. To use the default graph retrievers, you can ask questions as straightforward as:

```
query_engine = index.as_query_engine(include_text=True)

response = query_engine.query("What happened at Ryanair?")

print(str(response))
```

- QUICK

Here is where the defined LLM and embedding model come into play. Of course, you can also implement [custom retrievers](#) for potentially better accuracy.

- Partners

Summary

-

Constructing knowledge graphs without relying on LLMs is not only feasible but also cost-effective and efficient. By fine-tuning smaller, task-specific models, such as those in the Relik framework, you can achieve high-performance information extraction for your retrieval-augmented generation (RAG) applications.

[Contact Us](#)

Entity link^{ing}^{to} critical step in this process, ensures that recognized entities are accurately mapped to corresponding entries in a knowledge base, thereby maintaining the integrity and utility of the knowledge graph.

- By using frameworks like Relik and platforms such as Neo4j, it's possible to construct advanced knowledge graphs that facilitate complex data analysis and retrieval tasks, all without the high costs typically associated with deploying LLMs.
- This method not only makes powerful data processing tools more accessible but also promotes innovation and efficiency in information extraction workflows.

Make sure to give the [Relik library](#) a star. The code is available on [GitHub](#).

-

Solution

Partners

Entity Linking and Relationship Extraction With Relik in LlamaIndex was originally published in [Neo4j Developer Blog](#) on Medium, where people are continuing the conversation by highlighting and responding to this story.

OEM
entitygraphllamaindexllmNeo4j
Partners

aura DB

Technology
Neo4j's fully managed cloud
service



GraphAcademy
Academy
Us

Free online courses & certifications.
Join 100K+ Neo4j experts!

Newroom



Start Learning
Graphs4Good

Careers
Online Developer Conference
Happening November 7, 2024



Support
Aura
Save the Date

-
-
-

[Neo4j Community Disclaimer](#)

Contact Us



Author

Tomaž Bratanič, Graph ML and GenAI Research, Neo4j

Tomaž Bratanič works at the intersection of graphs and machine learning and generative AI.

[Email me blog updates!](#)

The information you provide will be used in accordance with the terms of our [privacy policy](#).

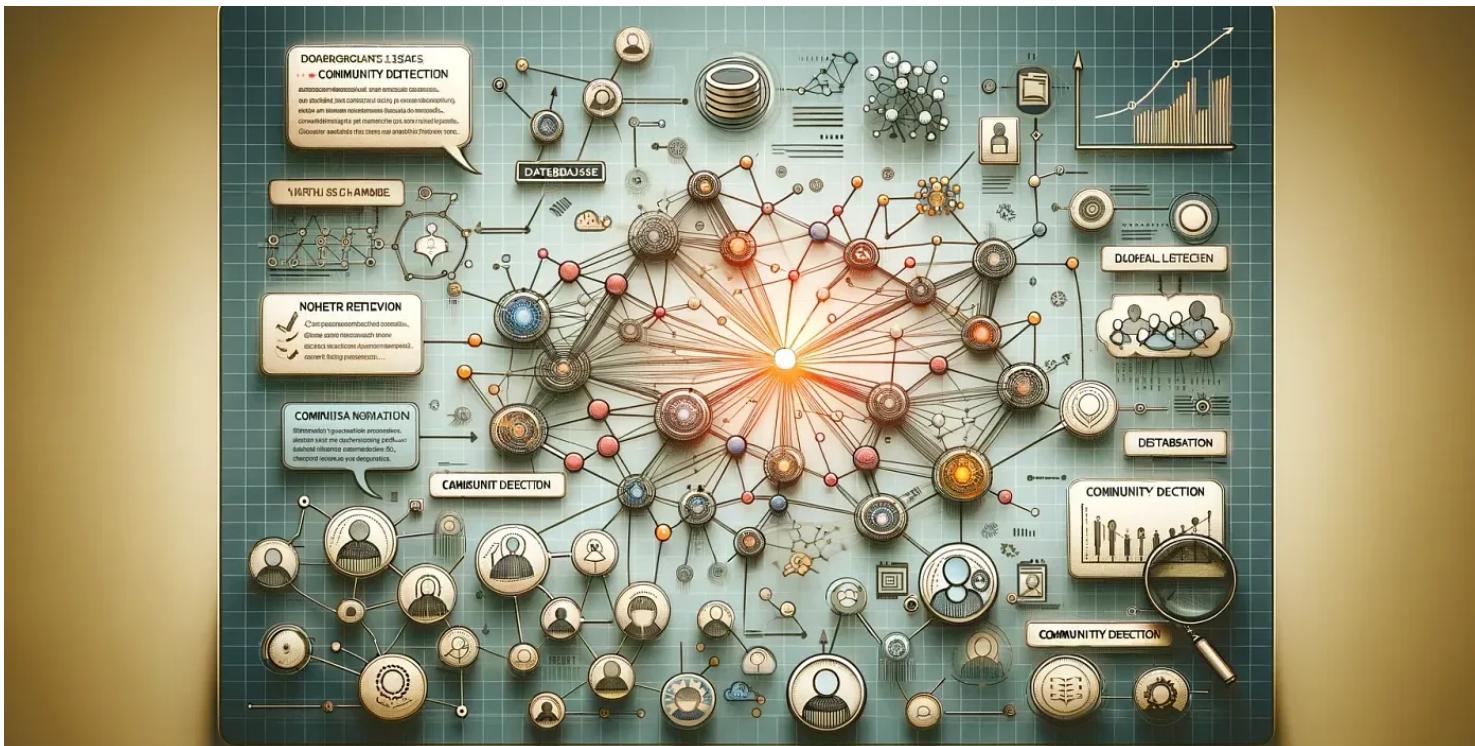
Related Articles

[Contact Us](#)



[Create a Neo4j GraphRAG Workflow Using LangChain and LangGraph](#) Aug 16 12 mins read

Contact Us



[Integrating Microsoft GraphRAG into Neo4j](#) Aug 13 16 mins read



[Build a Knowledge Graph-based Agent With Llama 3.1, NVIDIA NIM, and LangChain](#) Aug 08 5 mins read



Ready To Get Started?

Neo4j has been downloaded over 2 million times and has a large global community of developers.

[Get Started for Free](#)

Products

- [Neo4j Graph Database](#)
- [Neo4j AuraDB](#)
- [Neo4j Graph Data Science](#)
- [Deployment Center](#)
- [Professional Services](#)
- [Pricing](#)

Graph Tools

- [Neo4j Developer Tools](#)
- [Neo4j Workspace](#)
- [Neo4j Bloom](#)
- [Neo4j GraphQL Library](#)
- [Neo4j Data Connectors](#)
- [Cypher Query Language](#)

Use Cases

- [Generative AI](#)
- [Knowledge Graphs](#)
- [Industries & Use Cases](#)
- [Case Studies](#)
- [Customers](#)

Developers

- [Developer Home](#)
- [Documentation](#)
- [Deployment Center](#)
- [Developer Blog](#)
- [Community](#)
- [Virtual Events](#)
- [GraphAcademy](#)
- [Release Notes](#)

Data Scientists

- [Graph Data Science Home](#)
- [Data Science Documentation](#)
- [Get Started with Graph Data Science](#)
- [Data Science Community](#)
- [GraphAcademy for Data Science](#)

Learn

- [Resource Library](#)

Contact Us

- [Neo4j Blog](#)
- [GraphAcademy](#)
- [Research Center](#)
- [Case Studies](#)
- [Executive Insights](#)
- [Events Calendar](#)
- [GraphSummit](#)
- [Connections](#)
- [Webinars](#)

Partners

- [Find a Partner](#)
- [Become a Partner](#)
- [Solution Partners](#)
- [OEM Partners](#)
- [Technology Partners](#)
- [Partner Portal Login](#)

Company

- [About Us](#)
- [Newsroom](#)
- [Awards and Honors](#)
- [Graphs4Good](#)
- [Careers](#)
- [Culture](#)
- [Diversity](#)
- [Leadership](#)
- [Support](#)
- [Trust Center](#)

Contact Us →

- US: [1-855-636-4532](#)
- Sweden: [+46 171 480 113](#)
- UK: [+44 20 3868 3223](#)
- France: [+33 \(0\) 1 88 46 13 20](#)

Social Networks

© 2024 Neo4j, Inc.

[Terms](#) | [Privacy Policy](#) | [Sitemap](#)

[Anti-Corruption Policy](#)

Neo4j®, Neo Technology®, Cypher®, Neo4j® Bloom™, Neo4j® AuraDS™ and Neo4j® AuraDB™ are registered trademarks of Neo4j, Inc. All other marks are owned by their respective companies.

Contact Us